

# Introduction to SQL Joins

SQL **joins** allow you to retrieve data from multiple related tables based on a common column. In relational databases, data is often stored in **separate tables** to maintain normalization, so joins are essential for retrieving meaningful information.

There are several types of SQL joins:

1. **INNER JOIN** – Returns only matching records.
2. **LEFT JOIN (LEFT OUTER JOIN)** – Returns all records from the left table and matching ones from the right.
3. **CROSS JOIN** – Produces a Cartesian product of two tables.
4. **SELF JOIN** – Joins a table to itself.

**Note:** SQLite **does not support RIGHT JOIN or FULL OUTER JOIN** directly.

---

## Sample Database Setup

Before diving into joins, let's create and populate **sample tables**.

### Creating Sample Tables

```
PRAGMA foreign_keys = ON;

DROP TABLE IF EXISTS Employees;
DROP TABLE IF EXISTS Departments;

CREATE TABLE Departments (
    dept_id INTEGER PRIMARY KEY,
    dept_name TEXT UNIQUE NOT NULL
);

CREATE TABLE Employees (
    emp_id INTEGER PRIMARY KEY,
    emp_name TEXT NOT NULL,
    department_id INTEGER,
    FOREIGN KEY (department_id) REFERENCES Departments(dept_id) ON DELETE SET
NULL
);
```

# Adding Foreign Keys in SQLite3

In SQLite, **foreign keys** ensure that the referenced value exists in the parent table before being inserted into the child table. However, **foreign key constraints are disabled by default in SQLite**, so they need to be enabled explicitly.

Before executing any queries, ensure that foreign keys are enabled:

## Explanation of Foreign Key Constraints

- **department\_id INTEGER** in the Employees table is a foreign key.
  - **FOREIGN KEY (department\_id) REFERENCES Departments(dept\_id)** ensures that any department\_id in Employees **must exist in Departments**.
  - **ON DELETE SET NULL** ensures that if a department is deleted, the department\_id in Employees is set to NULL instead of being deleted.
- 

## Inserting Data (with Foreign Key Constraints)

Now, let's insert **valid data** into both tables.

```
INSERT INTO Departments (dept_id, dept_name) VALUES
(1, 'HR'),
(2, 'IT'),
(3, 'Finance'),
(4, 'Marketing');

INSERT INTO Employees (emp_id, emp_name, department_id) VALUES
(1, 'Alice', 1),      -- HR
(2, 'Bob', 2),        -- IT
(3, 'Charlie', 1),    -- HR
(4, 'David', NULL),   -- No department
(5, 'Eve', 3);        -- Finance
```

## Enforcing Referential Integrity

Now, let's **test** what happens if we try to insert an invalid department ID.

```
INSERT INTO Employees (emp_id, emp_name, department_id) VALUES
(6, 'Frank', 99);
```

This will fail because dept\_id 99 does not exist

**Error:** FOREIGN KEY constraint failed

SQLite **prevents** inserting department\_id = 99 because **dept\_id 99 does not exist in Departments**.

## Testing ON DELETE Behaviour

### Scenario: Deleting a Department

```
DELETE FROM Departments WHERE dept_id = 1;
```

Since we used **ON DELETE SET NULL**, employees **Alice and Charlie** will now have `department_id = NULL`, rather than being deleted.

### Verifying the Change

```
SELECT * FROM Employees;
```

### Updated Employees Table

emp_id	emp_name	department_id
1	Alice	NULL
2	Bob	2
3	Charlie	NULL
4	David	NULL
5	Eve	3

## SQL Join Types

Now, let's explore each type of join.

### 1. INNER JOIN

An **INNER JOIN** returns only records where there is a **match** in both tables.

#### Example Query

```
SELECT e.emp_id, e.emp_name, d.dept_name
FROM Employees e
INNER JOIN Departments d
ON e.department_id = d.dept_id;
```

### Output

emp_id	emp_name	dept_name
1	Alice	HR
2	Bob	IT
3	Charlie	HR

### Explanation

- Employees **Alice, Bob, and Charlie** appear because they have matching department IDs.
- **David (4)** is missing because his department\_id is NULL.
- **Marketing (dept\_id 4)** is missing because no employee belongs to it.

## 2. LEFT JOIN (LEFT OUTER JOIN)

A **LEFT JOIN** returns all records from the **left table** (Employees) and matching records from the right table (Departments). If no match is found, NULL values are returned.

### Example Query

```
SELECT e.emp_id, e.emp_name, d.dept_name
FROM Employees e
LEFT JOIN Departments d
ON e.department_id = d.dept_id;
```

### Output

emp_id	emp_name	dept_name
1	Alice	HR
2	Bob	IT
3	Charlie	HR
4	David	NULL
5	Eve	Finance

### Explanation

- All **employees** are included.
- **David** has no department, so dept\_name is NULL.

### 3. RIGHT JOIN (RIGHT OUTER JOIN)

A **RIGHT JOIN** returns all records from the **right table** (Departments) and matching records from the left table (Employees).

#### Example Query

```
-- Right Join
SELECT e.emp_id, e.emp_name, d.dept_name
FROM Employees e
RIGHT JOIN Departments d
ON e.department_id = d.dept_id;
```

#### Output

emp_id	emp_name	dept_name
1	Alice	HR
2	Bob	IT
3	Charlie	HR
5	Eve	Finance
NULL	NULL	Marketing

#### Explanation

The **Marketing department** is included, even though no employee is assigned to it.

**Note:** RIGHT JOIN is not supported in SQLite. Use LEFT JOIN by swapping table positions.

### 4 CROSS JOIN

A **CROSS JOIN** produces the **Cartesian Product**, meaning every row from the first table joins with every row from the second table.

#### Example Query

```
-- Cross Join
SELECT e.emp_name, d.dept_name
FROM Employees e
CROSS JOIN Departments d;
```

#### Output (if 5 employees and 4 departments)

Total  $5 \times 4 = 20$  rows.